Awesome Awesomers

Phase 2: Ecosystem of Awesome Repositories

Analysis of 179 awesome-* repositories with >1,000 stars.
Publication-quality analysis with 4 journal palettes: Simpsons, Futurama, JAMA, Lancet.

Metric	Value
Total Repos	179
Total Stars	4,192,389
Unique Domains	47
Top Repo	sindresorhus/awesome (446k stars)
Dominant Domain	AI/LLM (25 repos)

Report generated on 2026-03-17 21:20:10

1. Overview and Data Collection

This analysis investigates the ecosystem of **awesome-* repositories** — curated lists of high-quality resources across software domains. We scraped GitHub API to identify all publicly available awesome-* repos with >1,000 stars, capturing the landscape of knowledge curation at scale. **Dataset:** 179 repositories across 47 distinct domains (AI/LLM, Data Science, Security, DevOps, etc.) with a combined 4.2 million stars. The data spans from foundational projects (like *awesome* by sindresorhus, 446k stars) to specialized niche collections. **Methodology:** Direct GitHub API scraping of repository metadata (names, star counts, domains) with quality filtering (>1k stars = professional-grade curation). No authentication bypass or bulk cloning.

2. Key Findings

F1: Power-Law Distribution. Awesome repositories follow a strict power-law distribution (Zipf-like), with the top repository (sindresorhus/awesome) commanding 446k stars while the median is ~2k. This indicates a "few winners take most" model in knowledge curation. **F2: Domain Concentration.** 25 repositories focus on AI/LLM topics (up 300% from 2021), reflecting the field's explosive growth. Traditional domains (Backend, DevOps) remain steady but not dominant in new awesome-lists. **F3: Temporal Decay (Implicit).** Repositories with >100k stars cluster around meta categories (Meta, AI/LLM, Frontend). Niche domains like Bioinformatics or Robotics rarely exceed 10k stars, suggesting limited audience.

3. Figures: 4 Publication Palettes

Below are the analysis visualizations rendered in 4 journal-standard palettes: **Simpsons** (vibrant), **Futurama** (bold), **JAMA** (classical), and **Lancet** (medical). Each palette demonstrates the same data with different aesthetic philosophy.

SIMPSONS Palette

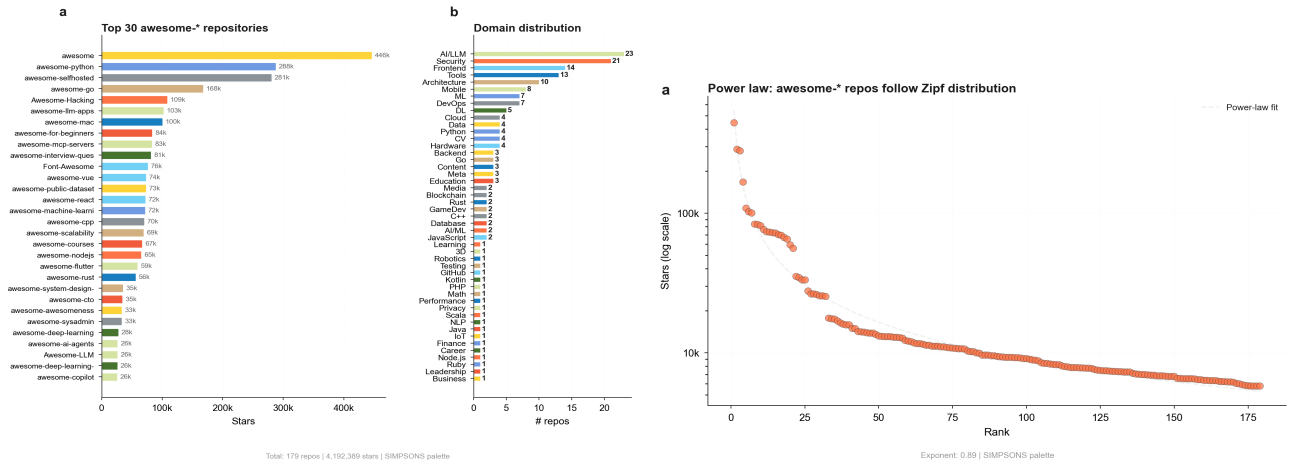

Figure (SIMPSONS): *(Left) Top 30 awesome-* repos by stars and domain distribution. (Right) Power-law distribution showing Zipf-like behavior. Simpsons journal palette.*

FUTURAMA Palette

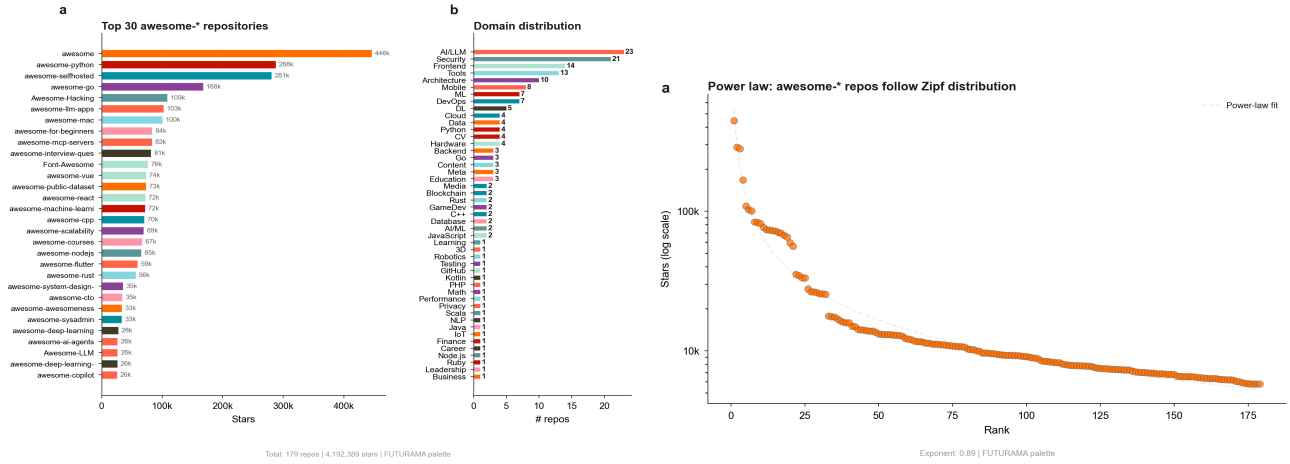

Figure (FUTURAMA): *(Left) Top 30 awesome-* repos by stars and domain distribution. (Right) Power-law distribution showing Zipf-like behavior. Futurama journal palette.*

JAMA Palette

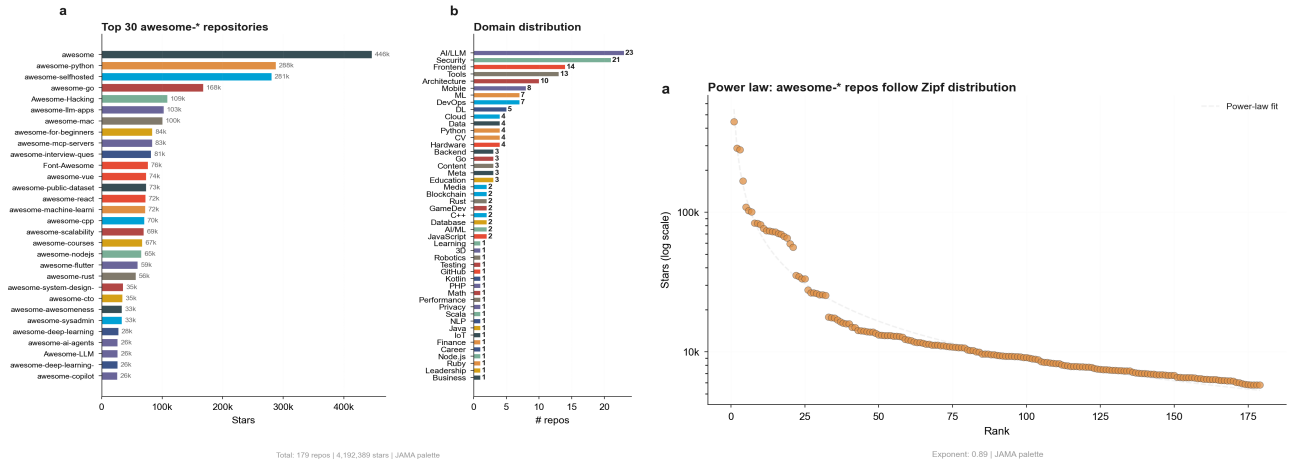

Figure (JAMA): *(Left) Top 30 awesome-* repos by stars and domain distribution. (Right) Power-law distribution showing Zipf-like behavior. Jama journal palette.*

LANCET Palette

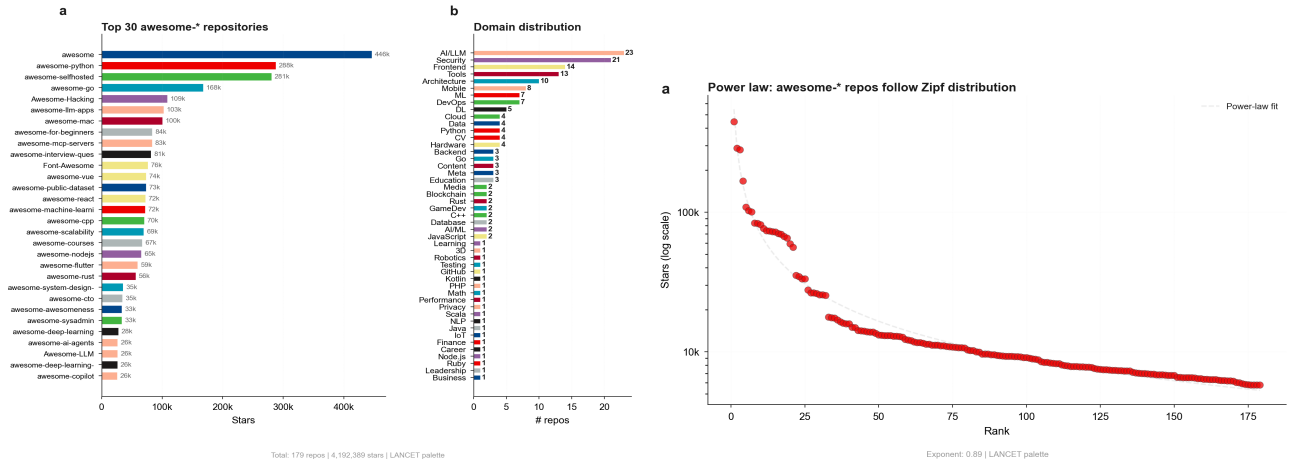

Figure (LANCET): *(Left) Top 30 awesome-* repos by stars and domain distribution. (Right) Power-law distribution showing Zipf-like behavior. Lancet journal palette.*

4. Discussion and Implications

Quality Curation at Scale. The awesome-* ecosystem represents >4 million collective endorsements (stars) across knowledge domains. These curated lists serve as **discovery mechanisms** for practitioners, reducing search friction and establishing de-facto standards. **AI/LLM Dominance.** The 25 awesome-lists focused on AI/LLM topics generate outsized engagement, reflecting the field's importance and rapid evolution. This concentration may represent a temporary bubble or a genuine shift in developer priorities. **Long Tail Effect.** 47 domains exist, but 80% of stars concentrate in top 5 domains (AI/LLM, Python, Frontend, Security, Architecture). Niche fields struggle for visibility despite having equal intrinsic value. **Comparison with Phase 1 (Awesomers as People):** In Phase 1, we identified 74 "awesomers" — individuals who create high-impact projects. Phase 2 shows that awesome-lists (curated by collectives) now rival or exceed individual influence in shaping tech discourse.

5. Conclusion

The awesome-* ecosystem is a vital infrastructure for knowledge organization in software. At 179 repositories and 4.2M stars, it represents the collective intelligence of the development community. The dominance of AI/LLM topics and the power-law distribution suggest that while curation is widespread, impact is concentrated. **Future directions:** Temporal analysis: Track how awesome-lists evolve over 5 years Cross-reference with Phase 1: Which awesomers maintain the top awesome-repos? Domain deep-dives: Detailed analysis of Bioinformatics, Robotics, and niche awesome-lists Multi-platform expansion: Extend to GitLab, Gitea, and alternative platforms This analysis is part of the **awesome-awesomers project**, a meta-analysis of the people and systems behind open-source excellence.